

Activity_Course 2 Automatidata project lab

December 13, 2023

1 Automatidata project

Course 2 - Get Started with Python

Welcome to the Automatidata Project!

You have just started as a data professional in a fictional data consulting firm, Automatidata. Their client, the New York City Taxi and Limousine Commission (New York City TLC), has hired the Automatidata team for its reputation in helping their clients develop data-based solutions.

The team is still in the early stages of the project. Previously, you were asked to complete a project proposal by your supervisor, DeShawn Washington. You have received notice that your project proposal has been approved and that New York City TLC has given the Automatidata team access to their data. To get clear insights, New York City TLC's data must be analyzed, key variables identified, and the dataset ensured it is ready for analysis.

A notebook was structured and prepared to help you in this project. Please complete the following questions.

2 Course 2 End-of-course project: Inspect and analyze data

In this activity, you will examine data provided and prepare it for analysis. This activity will help ensure the information is,

1. Ready to answer questions and yield insights
2. Ready for visualizations
3. Ready for future hypothesis testing and statistical methods

The purpose of this project is to investigate and understand the data provided.

The goal is to use a dataframe constructed within Python, perform a cursory inspection of the provided dataset, and inform team members of your findings.

This activity has three parts:

Part 1: Understand the situation * Prepare to understand and organize the provided taxi cab dataset and information.

Part 2: Understand the data

- Create a pandas dataframe for data learning, future exploratory data analysis (EDA), and statistical activities.
- Compile summary information about the data to inform next steps.

Part 3: Understand the variables

- Use insights from your examination of the summary data to guide deeper investigation into specific variables.

Follow the instructions and answer the following questions to complete the activity. Then, you will complete an Executive Summary using the questions listed on the PACE Strategy Document.

Be sure to complete this activity before moving on. The next course item will provide you with a completed exemplar to compare to your own work.

3 Identify data types and relevant variables using Python

4 PACE stages

Throughout these project notebooks, you'll see references to the problem-solving framework PACE. The following notebook components are labeled with the respective PACE stage: Plan, Analyze, Construct, and Execute.

4.1 PACE: Plan

Consider the questions in your PACE Strategy Document and those below to craft your response:

4.1.1 Task 1. Understand the situation

- How can you best prepare to understand and organize the provided taxi cab information?

Understand the current situation. Review the dataset to find out what data is in the dataset. We need to load the data and then prepare, clean data for analysis.

4.2 PACE: Analyze

Consider the questions in your PACE Strategy Document to reflect on the Analyze stage.

4.2.1 Task 2a. Build dataframe

Create a pandas dataframe for data learning, and future exploratory data analysis (EDA) and statistical activities.

Code the following,

- `import pandas as pd.` pandas is used for buidling dataframes.

- import numpy as np. numpy is imported with pandas
- df = pd.read_csv('Datasets\NYC taxi data.csv')

Note: pair the data object name `df` with pandas functions to manipulate data, such as `df.groupby()`.

Note: As shown in this cell, the dataset has been automatically loaded in for you. You do not need to download the .csv file, or provide more code, in order to access the dataset and proceed with this lab. Please continue with this activity by completing the following instructions.

```
[3]: #Import libraries and packages listed above
    ### YOUR CODE HERE ###
    import pandas as pd
    import numpy as np

    # Load dataset into dataframe
    df = pd.read_csv('2017_Yellow_Taxi_Trip_Data.csv')
    print("done")
```

done

4.2.2 Task 2b. Understand the data - Inspect the data

View and inspect summary information about the dataframe by coding the following:

1. `df.head(10)`
2. `df.info()`
3. `df.describe()`

Consider the following two questions:

Question 1: When reviewing the `df.info()` output, what do you notice about the different variables? Are there any null values? Are all of the variables numeric? Does anything else stand out?

Question 2: When reviewing the `df.describe()` output, what do you notice about the distributions of each variable? Are there any questionable values?

==> ENTER YOUR RESPONSE TO QUESTIONS 1 & 2 HERE

Answer 1: There're 3 objects, 8 float, and 7 int64 datatypes. There're 22699 records and 18 columns. No null values. 2 variables are date/time, and 1 variable is string. `column[0]` is `Unnamed:0`. Not sure what this column is about.

Answer 2: All variables are 22699 as count. There're negative numbers in `_fare_amount`, `extra_mta_tax`, `improvement_surcharge` and `total_amount`. There's 0 passenger count too. The maximum `_fare_amount` is a much larger value(\$1000) than the 25%-75% range of values. The maximum trip distance is over 33 miles.

```
[5]: #==> ENTER YOUR CODE HERE
    df.head(10)
```

```

[5]: Unnamed: 0 VendorID tpep_pickup_datetime tpep_dropoff_datetime \
0 24870114 2 03/25/2017 8:55:43 AM 03/25/2017 9:09:47 AM
1 35634249 1 04/11/2017 2:53:28 PM 04/11/2017 3:19:58 PM
2 106203690 1 12/15/2017 7:26:56 AM 12/15/2017 7:34:08 AM
3 38942136 2 05/07/2017 1:17:59 PM 05/07/2017 1:48:14 PM
4 30841670 2 04/15/2017 11:32:20 PM 04/15/2017 11:49:03 PM
5 23345809 2 03/25/2017 8:34:11 PM 03/25/2017 8:42:11 PM
6 37660487 2 05/03/2017 7:04:09 PM 05/03/2017 8:03:47 PM
7 69059411 2 08/15/2017 5:41:06 PM 08/15/2017 6:03:05 PM
8 8433159 2 02/04/2017 4:17:07 PM 02/04/2017 4:29:14 PM
9 95294817 1 11/10/2017 3:20:29 PM 11/10/2017 3:40:55 PM

```

```

passenger_count trip_distance RatecodeID store_and_fwd_flag \
0 6 3.34 1 N
1 1 1.80 1 N
2 1 1.00 1 N
3 1 3.70 1 N
4 1 4.37 1 N
5 6 2.30 1 N
6 1 12.83 1 N
7 1 2.98 1 N
8 1 1.20 1 N
9 1 1.60 1 N

```

```

PULocationID DOLocationID payment_type fare_amount extra mta_tax \
0 100 231 1 13.0 0.0 0.5
1 186 43 1 16.0 0.0 0.5
2 262 236 1 6.5 0.0 0.5
3 188 97 1 20.5 0.0 0.5
4 4 112 2 16.5 0.5 0.5
5 161 236 1 9.0 0.5 0.5
6 79 241 1 47.5 1.0 0.5
7 237 114 1 16.0 1.0 0.5
8 234 249 2 9.0 0.0 0.5
9 239 237 1 13.0 0.0 0.5

```

```

tip_amount tolls_amount improvement_surcharge total_amount
0 2.76 0.0 0.3 16.56
1 4.00 0.0 0.3 20.80
2 1.45 0.0 0.3 8.75
3 6.39 0.0 0.3 27.69
4 0.00 0.0 0.3 17.80
5 2.06 0.0 0.3 12.36
6 9.86 0.0 0.3 59.16
7 1.78 0.0 0.3 19.58
8 0.00 0.0 0.3 9.80
9 2.75 0.0 0.3 16.55

```

```
[3]: #==> ENTER YOUR CODE HERE
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 22699 entries, 0 to 22698
Data columns (total 18 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Unnamed: 0                            22699 non-null  int64
1   VendorID                             22699 non-null  int64
2   tpep_pickup_datetime                 22699 non-null  object
3   tpep_dropoff_datetime                22699 non-null  object
4   passenger_count                      22699 non-null  int64
5   trip_distance                       22699 non-null  float64
6   RatecodeID                          22699 non-null  int64
7   store_and_fwd_flag                   22699 non-null  object
8   PULocationID                        22699 non-null  int64
9   DOLocationID                        22699 non-null  int64
10  payment_type                         22699 non-null  int64
11  fare_amount                         22699 non-null  float64
12  extra                              22699 non-null  float64
13  mta_tax                            22699 non-null  float64
14  tip_amount                         22699 non-null  float64
15  tolls_amount                       22699 non-null  float64
16  improvement_surcharge               22699 non-null  float64
17  total_amount                       22699 non-null  float64
dtypes: float64(8), int64(7), object(3)
memory usage: 3.1+ MB
```

```
[4]: #==> ENTER YOUR CODE HERE
df.describe()
```

```
[4]:
```

	Unnamed: 0	VendorID	passenger_count	trip_distance	\
count	2.269900e+04	22699.000000	22699.000000	22699.000000	
mean	5.675849e+07	1.556236	1.642319	2.913313	
std	3.274493e+07	0.496838	1.285231	3.653171	
min	1.212700e+04	1.000000	0.000000	0.000000	
25%	2.852056e+07	1.000000	1.000000	0.990000	
50%	5.673150e+07	2.000000	1.000000	1.610000	
75%	8.537452e+07	2.000000	2.000000	3.060000	
max	1.134863e+08	2.000000	6.000000	33.960000	

	RatecodeID	PULocationID	DOLocationID	payment_type	fare_amount	\
count	22699.000000	22699.000000	22699.000000	22699.000000	22699.000000	
mean	1.043394	162.412353	161.527997	1.336887	13.026629	
std	0.708391	66.633373	70.139691	0.496211	13.243791	
min	1.000000	1.000000	1.000000	1.000000	-120.000000	

25%	1.000000	114.000000	112.000000	1.000000	6.500000
50%	1.000000	162.000000	162.000000	1.000000	9.500000
75%	1.000000	233.000000	233.000000	2.000000	14.500000
max	99.000000	265.000000	265.000000	4.000000	999.990000

	extra	mta_tax	tip_amount	tolls_amount	\
count	22699.000000	22699.000000	22699.000000	22699.000000	
mean	0.333275	0.497445	1.835781	0.312542	
std	0.463097	0.039465	2.800626	1.399212	
min	-1.000000	-0.500000	0.000000	0.000000	
25%	0.000000	0.500000	0.000000	0.000000	
50%	0.000000	0.500000	1.350000	0.000000	
75%	0.500000	0.500000	2.450000	0.000000	
max	4.500000	0.500000	200.000000	19.100000	

	improvement_surcharge	total_amount
count	22699.000000	22699.000000
mean	0.299551	16.310502
std	0.015673	16.097295
min	-0.300000	-120.300000
25%	0.300000	8.750000
50%	0.300000	11.800000
75%	0.300000	17.800000
max	0.300000	1200.290000

4.2.3 Task 2c. Understand the data - Investigate the variables

Sort and interpret the data table for two variables: `trip_distance` and `total_amount`.

Answer the following three questions:

Question 1: Sort your first variable (`trip_distance`) from maximum to minimum value, do the values seem normal?

Question 2: Sort by your second variable (`total_amount`), are any values unusual?

Question 3: Are the resulting rows similar for both sorts? Why or why not?

==> ENTER YOUR RESPONSES TO QUESTION 1-3 HERE

Answer 1: There're some 0s in `trip_distance`, and the maximum is over 33 miles, which is not normal.

Answer 2: There're some negative numbers in `total_amount`, and the maximum value is significantly higher than the others.

Answer 3: The resulting rows are not similar for both sorts. The most expensive rides are not necessarily the longest ones.

[4]: ## ==> ENTER YOUR CODE HERE

```
# Sort the data by trip distance from maximum to minimum value
df_sort=df.sort_values('trip_distance',ascending = False)
df_sort.head(10)
```

```
[4]:      Unnamed: 0  VendorID  tpep_pickup_datetime  tpep_dropoff_datetime  \
9280      51810714         2  06/18/2017 11:33:25 PM  06/19/2017 12:12:38 AM
13861     40523668         2   05/19/2017 8:20:21 AM   05/19/2017 9:20:30 AM
6064      49894023         2  06/13/2017 12:30:22 PM   06/13/2017 1:37:51 PM
10291     76319330         2  09/11/2017 11:41:04 AM   09/11/2017 12:18:58 PM
29         94052446         2  11/06/2017 8:30:50 PM  11/07/2017 12:00:00 AM
18130     90375786         1  10/26/2017 2:45:01 PM   10/26/2017 4:12:49 PM
5792      68023798         2   08/11/2017 2:14:01 PM   08/11/2017 3:17:31 PM
15350     77309977         2   09/14/2017 1:44:44 PM   09/14/2017 2:34:29 PM
10302     43431843         1   05/15/2017 8:11:34 AM   05/15/2017 9:03:16 AM
2592      51094874         2   06/16/2017 6:51:20 PM   06/16/2017 7:41:42 PM
```

```
      passenger_count  trip_distance  RatecodeID  store_and_fwd_flag  \
9280                 2          33.96           5                 N
13861                1          33.92           5                 N
6064                 1          32.72           3                 N
10291                1          31.95           4                 N
29                   1          30.83           1                 N
18130                1          30.50           1                 N
5792                 1          30.33           2                 N
15350                1          28.23           2                 N
10302                1          28.20           2                 N
2592                 1          27.97           2                 N
```

```
      PULocationID  DOLocationID  payment_type  fare_amount  extra  mta_tax  \
9280             132           265           2       150.00    0.0    0.0
13861            229           265           1       200.01    0.0    0.5
6064             138            1           1       107.00    0.0    0.0
10291            138           265           2       131.00    0.0    0.5
29              132            23           1        80.00    0.5    0.5
18130            132           220           1        90.50    0.0    0.5
5792             132           158           1        52.00    0.0    0.5
15350             13           132           1        52.00    0.0    0.5
10302             90           132           1        52.00    0.0    0.5
2592            261           132           2        52.00    4.5    0.5
```

```
      tip_amount  tolls_amount  improvement_surcharge  total_amount
9280          0.00          0.00                 0.3       150.30
13861         51.64          5.76                 0.3       258.21
6064         55.50         16.26                 0.3       179.06
10291          0.00          0.00                 0.3       131.80
```

29	18.56	11.52	0.3	111.38
18130	19.85	8.16	0.3	119.31
5792	14.64	5.76	0.3	73.20
15350	4.40	5.76	0.3	62.96
10302	11.71	5.76	0.3	70.27
2592	0.00	5.76	0.3	63.06

[7]: *#==> ENTER YOUR CODE HERE*

```
# Sort the data by total amount and print the top 20 values
total_amount_sorted=df.
↳sort_values('total_amount',ascending=False)['total_amount']
total_amount_sorted.head(20)
```

```
[7]: 8476      1200.29
      20312      450.30
      13861      258.21
      12511      233.74
      15474      211.80
      6064      179.06
      16379      157.06
      3582      152.30
      11269      151.82
      9280      150.30
      1928      137.80
      10291      131.80
      6708      126.00
      11608      123.30
      908       121.56
      7281      120.96
      18130      119.31
      13621      115.94
      13359      111.95
      29        111.38
Name: total_amount, dtype: float64
```

[10]: *#==> ENTER YOUR CODE HERE*

```
# Sort the data by total amount and print the bottom 20 values
total_amount_sorted.tail(20)
```

```
[10]: 14283      0.31
      19067      0.30
      10506      0.00
      5722      0.00
      4402      0.00
      22566      0.00
```



```

1646      -3.30
18565     -3.80
314       -3.80
5758      -3.80
5448      -4.30
4423      -4.30
10281     -4.30
8204      -4.80
20317     -4.80
11204     -5.30
14714     -5.30
17602     -5.80
20698     -5.80
12944    -120.30
Name: total_amount, dtype: float64

```

```

[14]: # show trip_distance, fare_amount, total_amount, order by total_amount
      ↪descending.
df.sort_values('total_amount',ascending=False).iloc[:20,[5,11,17]]

```

```

[14]:
      trip_distance  fare_amount  total_amount
8476             2.60        999.99        1200.29
20312             0.00        450.00         450.30
13861            33.92        200.01         258.21
12511             0.00        175.00         233.74
15474             0.00        200.00         211.80
6064             32.72        107.00         179.06
16379            25.50        140.00         157.06
3582              7.30        152.00         152.30
11269             0.00        120.00         151.82
9280             33.96        150.00         150.30
1928             12.50        120.00         137.80
10291            31.95        131.00         131.80
6708              0.32        100.00         126.00
11608            23.00         99.50         123.30
908             26.12        100.00         121.56
7281             0.00        100.00         120.96
18130            30.50         90.50         119.31
13621            19.80        105.00         115.94
13359             0.00         75.00         111.95
29             30.83         80.00         111.38

```

```

[21]: #==> ENTER YOUR CODE HERE

# How many of each payment type are represented in the data?
df['payment_type'].value_counts()

```

```
[21]: 1    15265
      2     7267
      3     121
      4      46
      Name: payment_type, dtype: int64
```

According to the data dictionary, the payment method was encoded as follows:

1 = Credit card
2 = Cash
3 = No charge
4 = Dispute
5 = Unknown
6 = Voided trip

```
[15]: #==> ENTER YOUR CODE HERE

# What is the average tip for trips paid for with credit card?
mask=df['payment_type']==1
credit_average_tip=df[mask]['tip_amount'].mean()
print("Credit card payment type average tip amount: $",credit_average_tip)

#==> ENTER YOUR CODE HERE

# What is the average tip for trips paid for with cash?
mask=df['payment_type']==2
credit_average_tip=df[mask]['tip_amount'].mean()
print("Cash payment type average tip amount: $",credit_average_tip)
```

Credit card payment type average tip amount: \$ 2.7298001965279934
Cash payment type average tip amount: \$ 0.0

```
[26]: #==> ENTER YOUR CODE HERE

# How many times is each vendor ID represented in the data?
df['VendorID'].value_counts()
```

```
[26]: 2    12626
      1    10073
      Name: VendorID, dtype: int64
```

```
[18]: #==> ENTER YOUR CODE HERE

# What is the mean total amount for each vendor?
df.groupby('VendorID').mean()[['total_amount']]
```

```
[18]:          total_amount
VendorID
```

```
1          16.298119
2          16.320382
```

```
[20]: #==> ENTER YOUR CODE HERE

# Filter the data for credit card payments only
mask = df['payment_type']==1
credit_card = df[mask]

#==> ENTER YOUR CODE HERE

# Filter the credit-card-only data for passenger count only
credit_card['passenger_count'].value_counts()
```

```
[20]: 1    10977
      2     2168
      5      775
      3      600
      6      451
      4      267
      0       27
      Name: passenger_count, dtype: int64
```

```
[26]: #==> ENTER YOUR CODE HERE

# Calculate the average tip amount for each passenger count (credit card,
↳ payments only)
#average_tip = df[mask]['tip_amount']/df[mask]['passenger_count']
#average_tip
credit_card.groupby('passenger_count').mean(numeric_only=True)[['tip_amount']]
```

```
[26]:          tip_amount
passenger_count
0          2.610370
1          2.714681
2          2.829949
3          2.726800
4          2.607753
5          2.762645
6          2.643326
```

4.3 PACE: Construct

Note: The Construct stage does not apply to this workflow. The PACE framework can be adapted to fit the specific requirements of any project.

4.4 PACE: Execute

Consider the questions in your PACE Strategy Document and those below to craft your response.

4.4.1 Given your efforts, what can you summarize for DeShawn and the data team?

Note for Learners: Your notebook should contain data that can address Luana's requests. Which two variables are most helpful for building a predictive model for the client: NYC TLC?

==> ENTER YOUR RESPONSE HERE

In the dataset the column DTypes are 3 objects, 8 float, and 7 int64. There's no null values in the dataset. The relevant columns are VendorID, trip_distance, passenger_count, payment_type, fare_amount, total_amount, tip_amount. The irrelevant columns are Unnamed:0, store_and_fwd_flag, PULocationID, DOLocationID, extra, mta_tax, tolls_amount, improvement_surcharge. The most helpful variables are total_amount and trip_distance, because these variables show a picture of a taxi cab ride.

Congratulations! You've completed this lab. However, you may not notice a green check mark next to this item on Coursera's platform. Please continue your progress regardless of the check mark. Just click on the "save" icon at the top of this notebook to ensure your work has been logged.